

Rahul Gupta

■ alice.chen@email.com | ■ +1 (415) 555-0182 | ■ linkedin.com/in/alicechen | ■ San Francisco, CA

PROFESSIONAL SUMMARY

Accomplished Senior Data Scientist with 7 years of experience delivering high-impact machine learning solutions across NLP, computer vision, and predictive analytics domains. Proven track record of leading end-to-end ML pipelines from data ingestion to production deployment. Expert in Python, TensorFlow, PyTorch, and cloud-based MLOps practices. Published researcher with 3 peer-reviewed papers in applied machine learning.

TECHNICAL SKILLS

Languages: Python, R, SQL, Scala, Bash

ML/DL Frameworks: TensorFlow, PyTorch, Scikit-learn, XGBoost, LightGBM, Hugging Face Transformers

MLOps & Cloud: MLflow, Kubeflow, Docker, Kubernetes, AWS SageMaker, GCP Vertex AI, Azure ML

Data Engineering: Apache Spark, Kafka, Airflow, dbt, Pandas, NumPy

Databases: PostgreSQL, MongoDB, BigQuery, Redshift, Snowflake, Pinecone (Vector DB)

Visualization: Tableau, Power BI, Matplotlib, Seaborn, Plotly

NLP Tools: NLTK, spaCy, BERT, GPT fine-tuning, LangChain, RAG pipelines

Statistical Methods: A/B Testing, Bayesian Inference, Causal Inference, Time Series (ARIMA, Prophet)

PROFESSIONAL EXPERIENCE

Senior Data Scientist | Stripe Inc., San Francisco, CA

January 2021 – Present

- Led development of a real-time fraud detection ML system processing 2M+ transactions/day, reducing fraud losses by 34%.
- Built and deployed NLP pipeline for automated dispute classification using BERT fine-tuning, achieving 91% accuracy.
- Designed and implemented A/B testing framework adopted org-wide for model comparison and feature validation.
- Mentored 4 junior data scientists; established MLOps best practices using MLflow and Kubernetes.
- Collaborated with product and engineering teams to integrate ML models into production APIs (FastAPI + Docker).

Data Scientist | LinkedIn, Sunnyvale, CA

June 2018 – December 2020

- Developed recommendation system for job matching using collaborative filtering and graph neural networks.
- Built time series forecasting models (Prophet, LSTM) for talent supply-demand prediction with 87% accuracy.
- Reduced model training time by 60% by migrating pipelines to Apache Spark on Databricks.
- Performed statistical analysis and causal inference studies to measure platform feature impact.

Junior Data Scientist | Accenture Applied Intelligence, New York, NY

August 2016 – May 2018

- Built classification models for customer churn prediction (XGBoost, Random Forest) with AUC-ROC of 0.89.
- Conducted EDA and feature engineering for retail demand forecasting projects.
- Created interactive Tableau dashboards for C-suite reporting.

EDUCATION

M.S. in Computer Science (Machine Learning specialization)

Stanford University, Stanford, CA | GPA: 3.92/4.0 | 2016

B.S. in Statistics & Data Science

UC Berkeley, Berkeley, CA | GPA: 3.85/4.0 | 2014

CERTIFICATIONS & PUBLICATIONS

- AWS Certified Machine Learning – Specialty (2023)
- Google Professional Data Engineer (2022)
- Deep Learning Specialization – Coursera/DeepLearning.AI (2020)
- Published: "Scalable Fraud Detection with Graph Neural Networks" – NeurIPS 2022
- Published: "Causal Inference in Online Marketplace A/B Tests" – KDD 2021